

Strengthening Weak Links in the PDF Trust Chain

Mark Tullsen, William Harris, Peter Wyatt

May 26, 2022

Outline

- 1 Pre-DOM (Pre Document Object Model)
- 2 Modeling Pre-DOM: Highlights
- 3 Conclusion
- 4 Preview: A Next Step

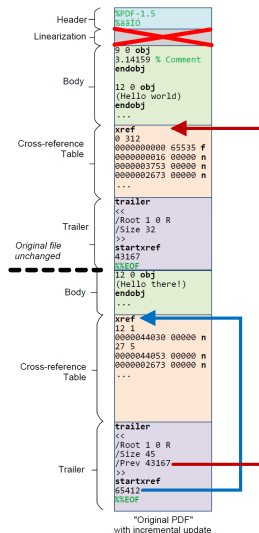
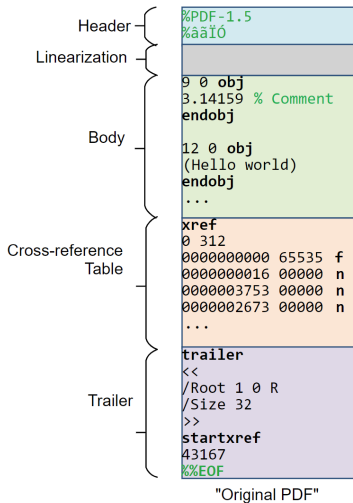
1 Pre-DOM (Pre Document Object Model)

2 Modeling Pre-DOM: Highlights

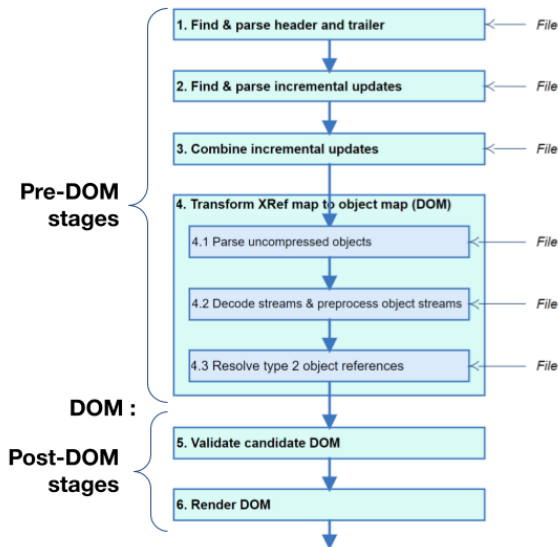
3 Conclusion

4 Preview: A Next Step

PDF Complexity?



PDF Trust Chain: Pre-DOM and Post-DOM



Vulnerabilities Occurring Primarily Pre-DOM

- Schizophrenic files (different tools, different renderings)
- Polyglot files (file being in 2+ formats)
- Shadow attacks
 - i.e., attacker can
 - add "shadow content" that is PDF-signed,
 - after signing, can update-at-will (revealing shadow content)
 - without giving clear warnings to user.
 - possible because of ability to sign *dead objects* and *cavities*
- Multiple places for hidden/unused/malicious data in PDF
 - non-obvious places, unnoticed when "simply parsing"
 - e.g., shadow-attacks
 - dead bytes, dead objects, dead updates, dead linearization sections, etc.

PDF, and Pre-DOM, Challenges

- Lack of formality in standard. Thus, implementations:
 - are more effort
 - over implement, under implement, wrongly implement
- No definition of acceptable, reasonable error recovery
- Less than ideal design that reflects 27 years of an evolving standard
- Pre-DOM processing
 - is where many parsing errors & recovery occur
 - is non-trivial
 - involves multiple interdependent features
 - involves multiple redundant features
 - schizophrenic if these features aren't mutually consistent


1 Pre-DOM (Pre Document Object Model)

2 Modeling Pre-DOM: Highlights

3 Conclusion

4 Preview: A Next Step

Stage 4 Sub-Stages: Transform XRef Map to Object Map




```
...
100 3 0 obj 99 endobj
123 % object 4 is not here
151 5 0 obj
<<
/Type /ObjStm
/Length 3 0 R % indirect!
/N 2 % 2 objects; (potentially indirect)
/First 10 % offset to 1st object (potentially indirect)
>>
```

Parser \neq Validator


- A surprising source of mis-communication.
- ...

A Validator (Parser \neq Validator)



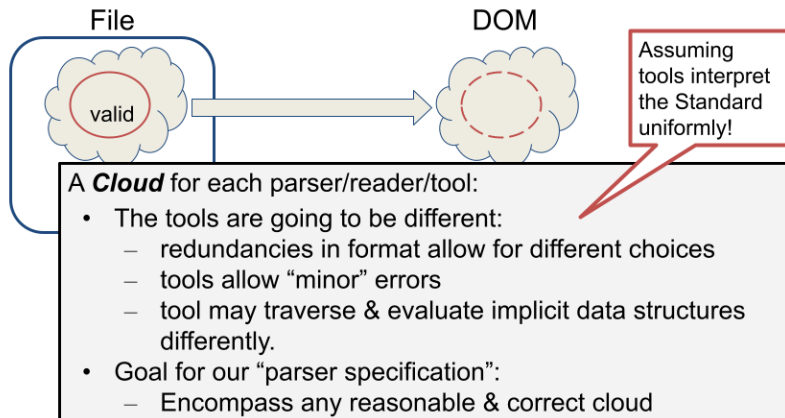
VALIDATOR: only valid PDFs can produce DOM (must Fail otherwise)

A Parser (Parser \neq Validator)



PARSER: efficiently, construct the correct DOM when a valid PDF

A Very Accepting Parser (Parser \neq Validator)



Turning Parser into Validator

Parser specification is designed to be

- understandable: clear, pure Haskell
- phased, clearly terminating (get parallelizability for free)
- very lazy "Parser" (big input cloud)

We can extend spec into a validator, orthogonally, via "validate" constructs (turning on/off on with command-line flag). E.g.,

do

```
(xrefRaw , xrefEndOff) <- pXrefRaw
```

```
validate $
```

```
    verifyXrefRaw xrefRaw — ensures no duplicates
```

1 Pre-DOM (Pre Document Object Model)

2 Modeling Pre-DOM: Highlights

3 Conclusion

4 Preview: A Next Step

Accomplishments

- A specification for pre-DOM parsing/computation
 - Clarifies some subtle issues in PDF Standard
 - A growing list of PDF Association “issues” that we have contributed to creating [23,24,...,30]
- Cause *and* effect of
 - unique tool for displaying updates & cavities

Future

- Not accomplished yet
 - the less interesting/subtle parts specified/implemented
 - integrated with our primitive, daedalus-generated parsers to create a tool.
- Create a full pre-DOM tool that
 - supports further PDF features (hybrids, compression, ...)
 - add support for commonly allowed “exuberances”
 - add more “validate”s to get closer to a *validator*.

1 Pre-DOM (Pre Document Object Model)

2 Modeling Pre-DOM: Highlights

3 Conclusion

4 Preview: A Next Step

Implementation?

Tools & renderers rarely need (*demand*) the whole PDF

- reading?
- parsing??
- semantic checks???

Thus, this

`parsePDF :: FileData → Maybe PDFAbstractSyntax`

is not going to be used in practice!

One Solution ...

- For complex formats,
 - tools are "projections": rarely used parse/validate all.
 - may have alternate "parsing paths" we want to take
 - e.g., metadata, page 1, text-only
- Shotgun Parsers?
 - ... the deadliest of patterns: "Input data checking, handling interspersed with processing logic"
- I.e., we provide multiple parsers where the following is interspersed through code and the relation between these is **not specified**:

```
parseA  :: Offset -> IO A
parseB  :: Offset -> IO B
parseC  :: Offset -> IO C
validateA :: A -> IO ()
validateB :: A -> B -> IO ()
```

Better Solution, Parser as API

We provide four inter-dependent calls (not *entry points*):

```
parseHdrTrlr  :: FileData → IO HdrTrlr
parseUpdates  :: HdrTrlr  → IO [Updates]
createXRef    :: [Updates] → IO XRef
derefObjId    :: ObjId    → XRef → IO PdfValue
```

(The returned types can be as abstract as we wish.)

Using this, we write abstractions on the above:

```
getInitialUpdate :: FileData → IO XRef
getRootValue      :: HdrTrailer → XRef → PdfValue
getPageTree      :: XRef → Tree PdfValue
```